

Intro to Heap

Nathan Huckleberry

University of Texas at Austin

April 22, 2021

Table of Contents

Heap Basics

Specifics of glibc

Attack Goals

Use After Free

Double Free

Advanced Tactics

Table of Contents

Heap Basics

Specifics of glibc

Attack Goals

Use After Free

Double Free

Advanced Tactics

What does a heap do

- ▶ We ask the heap for n bytes of memory
- ▶ It allocates n bytes of memory for us to use
- ▶ We eventually give the memory back when we're done

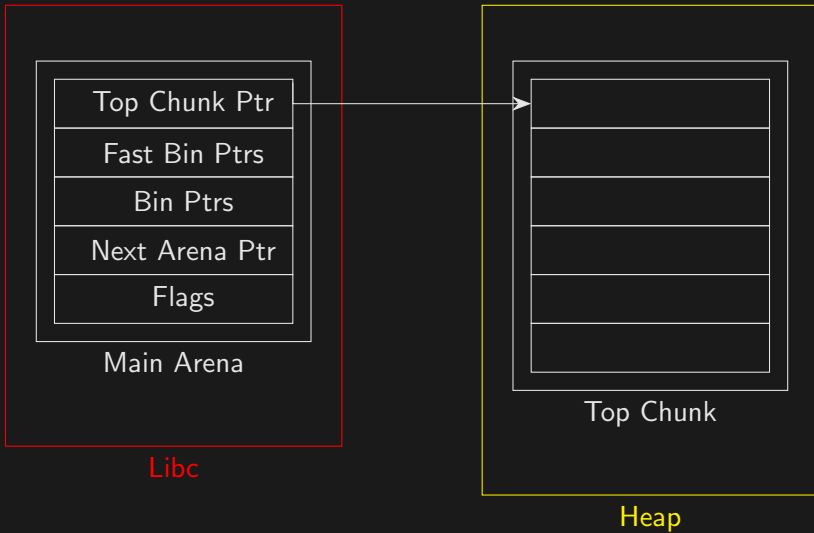
Dynamic Memory

- ▶ A program can only request entire pages (contiguous 4096 byte sections) from the kernel
- ▶ We want to break up these pages into smaller chunks
- ▶ Need some code that manages these chunks for us

Heap Initialization

- ▶ The heap will request a large chunk called the top chunk or wilderness chunk
- ▶ Subsequent requests from malloc will get a piece of the wilderness chunk
- ▶ The heap main arena keeps track of where the top-chunk lives

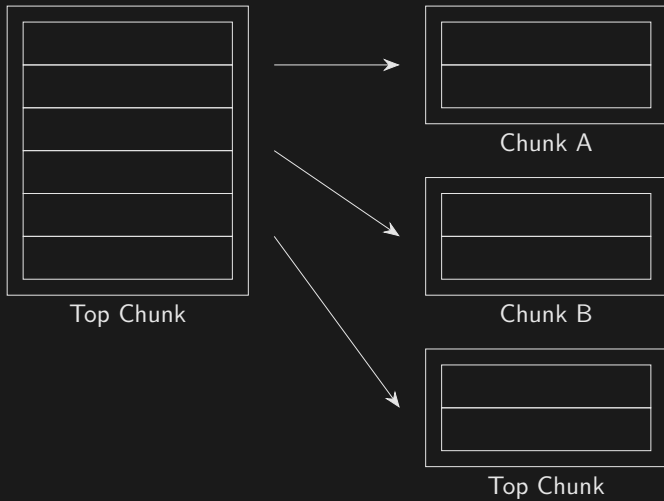
Wilderness Chunk



Allocation

- ▶ When the program requests memory, memory is cut off the top chunk
- ▶ The top chunk shrinks
- ▶ The memory allocated from the top chunk is returned to the program

Allocation




Freeing

- ▶ When we free memory, we want it to be made usable for future allocations
- ▶ We can only recombine with the top chunk if the freed chunk is next to the top chunk in memory
- ▶ Most cases we can't recombine

Freeing


- ▶ We can't recombine A and the top chunk since B is in the middle



Free Chunk A



Chunk B



Top Chunk

Freeing

- ▶ Instead we put it in a list of free chunks
- ▶ When we allocate later, we look if there are any suitably sized free chunks in the list before taking from the top chunk
- ▶ This free list lookup needs to be extremely efficient

Freeing



- ▶ An implementation like above may require enumeration of the entire list for large allocations
- ▶ Not efficient enough

Freeing Efficiency

- ▶ Chunks are bucketed based on size
- ▶ For each size chunk, there is a separate free list
- ▶ These are called bins

Fast Bins

- ▶ Bins for small sized chunks are called fast bins (0x10 - 0xb0)
- ▶ They have simplified handling since we expect small allocations to be freed/allocated very often
- ▶ Fast bins are the simplest for CTF challenges

Fast Bins

- ▶ Fast bins are a singly linked list of free chunks all of the same size
- ▶ Allocation is made for small chunk size
 - ▶ Round the allocation size up to nearest fastbin size
 - ▶ Return first element in fastbin

Fast Bins



- Every chunk in the list is the same size, so we don't need to search

Other Bin Types

- ▶ There are other types of bins, namely small bins, large bins and unsorted bins
- ▶ These bins are double linked, circular linked lists
- ▶ The head/tail are in the arena struct

Table of Contents

Heap Basics

Specifics of glibc

Attack Goals

Use After Free

Double Free

Advanced Tactics

Chunk Layout

- ▶ In glibc the pointer to the next chunk in the free list is called `fd`
- ▶ The pointer to the previous chunk in the free list is called `bk`
- ▶ Each chunk is allocated with an additional 8 bytes specifying the size

Chunk Layout

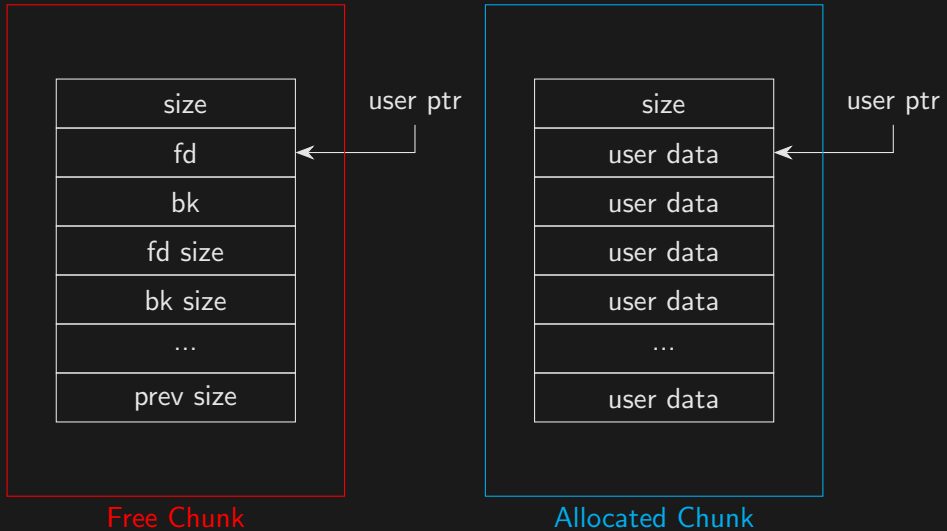


Table of Contents

Heap Basics

Specifics of glibc

Attack Goals

Use After Free

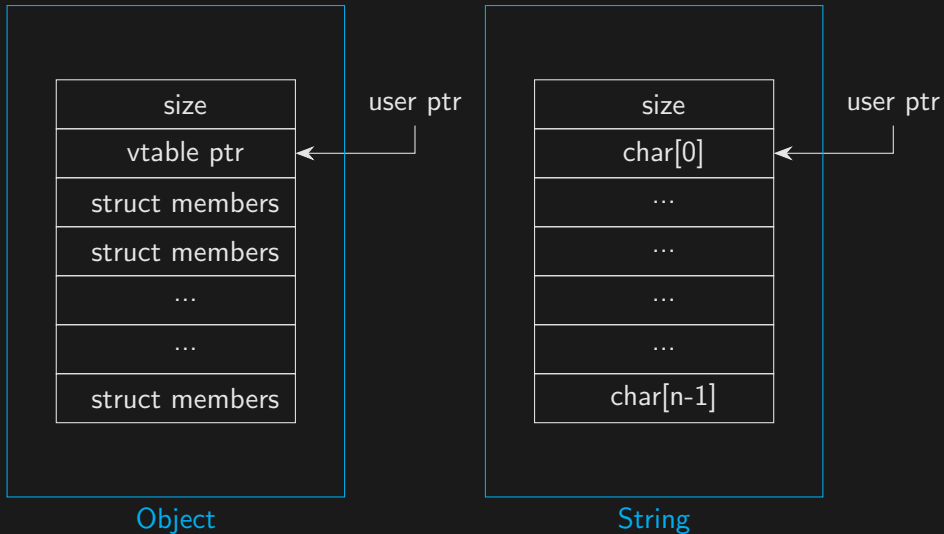
Double Free

Advanced Tactics

Overlapping Allocations

- ▶ Cause the heap to allocate two different structs at the same memory
- ▶ Use these allocations to corrupt the structs
 - ▶ Overwrite a function pointer in the struct
 - ▶ Cause a stack overflow by corrupting in a creative way

Overlapping Allocations



Forged Chunk

- ▶ Cause the heap to allocate a struct outside of the heap region
- ▶ Overwrite GOT entries or return address directly

Table of Contents

Heap Basics

Specifics of glibc

Attack Goals

Use After Free

Double Free

Advanced Tactics

Overlapping Allocation UAF

- ▶ We're allowed to write to the chunk *after* it's been freed
- ▶ The next time we call malloc with the chunk size, we get the same memory
- ▶ Overlapping allocation comes for free

Overlapping Allocation UAF



UAF Chunk A



Free Chunk A

- ▶ Allocate a chunk of size 0x80
- ▶ We get chunk A
- ▶ We now have two allocations of chunk A

Overlapping Allocation UAF



UAF Chunk A



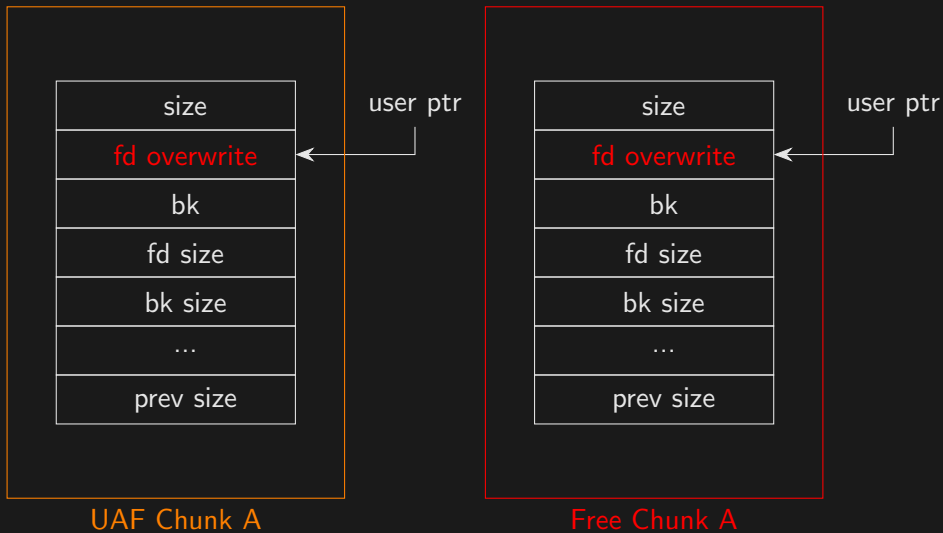
Allocated Chunk A

- ▶ Allocate a chunk of size 0x80
- ▶ We get chunk A
- ▶ We now have two allocations of chunk A


Forged Chunk UAF

- ▶ By writing to the freed chunk we can mess with linked list pointers

Forged Chunk UAF



Fast Bins UAF



- ▶ After allocating 3 chunks of size 0x80, the heap will return a chunk at an attacker controlled address
- ▶ We can then use this chunk to overwrite arbitrary memory (GOT, return address, malloc hook, etc)

Fast Bins UAF



- Overwriting GOT pointers will allow us to take control of the program

In Practice

- ▶ Malloc checks to make sure the size matches with the bin we're allocating from
- ▶ We need to write *size* somewhere in memory before we can allocate a forged chunk

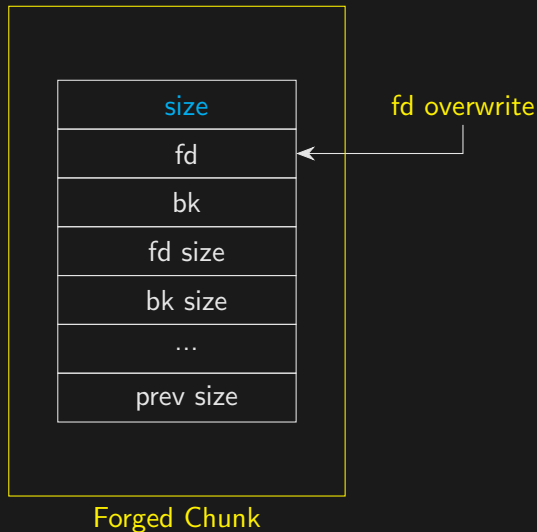


Table of Contents

Heap Basics

Specifics of glibc

Attack Goals

Use After Free


Double Free

Advanced Tactics

Double Free

- ▶ We're allowed to write to allocated chunks
- ▶ We're allowed to free a pointer twice

Double Free




- ▶ Freeing the same pointer twice creates a cycle in the free list
- ▶ Allocating will give us chunk A and the free list stays the same

Double Free



- ▶ Allocate a chunk of size 0x80
- ▶ We get free chunk A
- ▶ The head of the free list stays free chunk A
- ▶ This is enough for an overlapping allocation

Double Free



- ▶ We do an fd overwrite as in the UAF case
- ▶ Allocating again gives us free chunk A
- ▶ Allocating again gives us our forged chunk

In Practice

- ▶ The heap has checks to make sure that you're not freeing the head of the free list
- ▶ To get around this, we free an additional chunk in between

In Practice



- ▶ Free A, Free B, Free A again
- ▶ There's still a cycle, it's just length 2 now

Table of Contents

Heap Basics

Specifics of glibc

Attack Goals

Use After Free

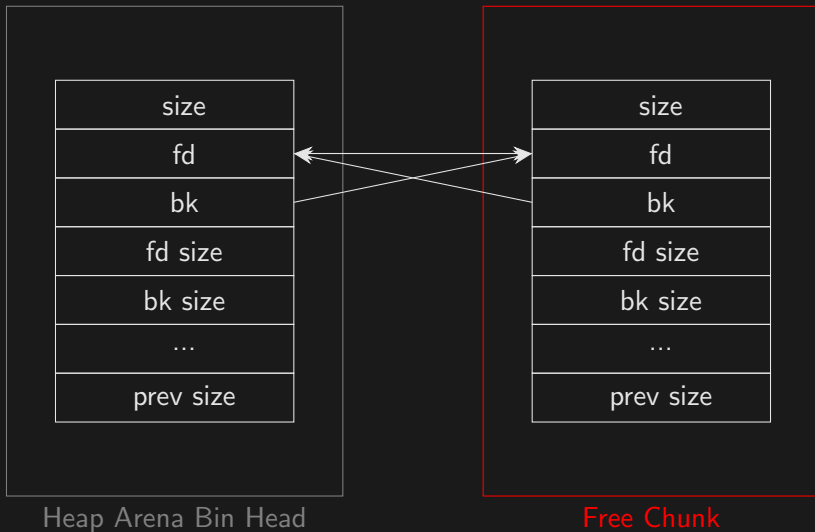
Double Free

Advanced Tactics

Libc Leaks with UAF

- ▶ For allocations larger than fast bins, the free list is doubly linked and circular
- ▶ We can use this to leak heap addresses
- ▶ We must be aware of the list head in the heap arena
 - ▶ A fake chunk in libc so the heap can keep track of the free list

Libc Leaks with UAF



Libc Leaks with UAF

- ▶ Printing the fd or bk pointers will give a leaked libc address
- ▶ Some address inside the main_arena struct